

tuning the toolkit's intake valves like a Stradivarius, coming right up. here's the retrofit that gives you:

- an incremental directory cache (skip re-ingesting files that haven't changed),
- a tiny HDF5 walker that dives one level into stubborn group layouts, and
- a `--merge-policy` switch to resolve "eigenpairs vs operator-derived eigenpairs" when shapes fight.

everything below drops cleanly into the structure we already set up.

---

## 1) Incremental cache over directories

**new file:**

**willowlab/ingest/cache\_dir.py**

```
# willowlab/ingest/cache_dir.py
```

```
import os, json, hashlib, pathlib, time
```

```
from typing import Dict, Any, Optional
```

```
DEFAULT_CACHE_NAME = "ingest_cache.json"
```

```
def _hash_file(path: str, block_size: int = 1024 * 1024) -> str:
```

```
    """SHA256 of file contents, streamed (works for big blobs)."""
```

```
    h = hashlib.sha256()
```

```
    with open(path, "rb") as f:
```

```
        while True:
```

```
            chunk = f.read(block_size)
```

```
        if not chunk: break

        h.update(chunk)

    return h.hexdigest()
```

```
def _mtime(path: str) -> float:
```

```
    try:

        return os.path.getmtime(path)

    except OSError:

        return 0.0
```

```
class DirCache:
```

```
    """
```

Directory-level cache:

- remembers file SHA256 and size/mtime
- provides 'should\_process' decision
- records arbitrary ingest metadata (e.g., where we saved the Willow bundle)

```
    """
```

```
def __init__(self, root_dir: str, cache_path: Optional[str] = None):
```

```
    self.root_dir = os.path.abspath(root_dir)
```

```
    self.cache_path = cache_path or os.path.join(self.root_dir, DEFAULT_CACHE_NAME)
```

```
    self._db: Dict[str, Any] = {}
```

```
    if os.path.exists(self.cache_path):
```

```
        try:
```

```
            self._db = json.load(open(self.cache_path, "r"))
```

```
except Exception:
```

```
    self._db = {}
```

```
def file_record(self, path: str) -> Optional[Dict[str, Any]]:
```

```
    key = os.path.relpath(os.path.abspath(path), self.root_dir)
```

```
    return self._db.get(key)
```

```
def should_process(self, path: str) -> bool:
```

```
    """Skip if same hash & size persisted; true means reprocess."""
```

```
    abspath = os.path.abspath(path)
```

```
    key = os.path.relpath(abspath, self.root_dir)
```

```
    size = os.path.getsize(abspath)
```

```
    mtime = _mtime(abspath)
```

```
    rec = self._db.get(key)
```

```
    if not rec:
```

```
        return True
```

```
    # Fast path: if size or mtime changed, re-hash and reprocess.
```

```
    if rec.get("size") != size or abs(rec.get("mtime", 0.0) - mtime) > 1e-6:
```

```
        return True
```

```
    # No change in size/mtime -> trust it as unchanged (cheap). If you want belt+braces,  
    re-hash here.
```

```
    return False
```

```
def update(self, path: str, **ingest_meta):
```

```
    abspath = os.path.abspath(path)
```

```

key = os.path.relpath(abspath, self.root_dir)

size = os.path.getsize(abspath)

mtime = _mtime(abspath)

# Hash only when we actually processed; keeps idle runs fast.

sha = _hash_file(abspath)

self._db[key] = {"size": size, "mtime": mtime, "sha256": sha, "ingest_meta": ingest_meta,
"updated": time.time()}

def persist(self):

    pathlib.Path(self.cache_path).parent.mkdir(parents=True, exist_ok=True)

    with open(self.cache_path, "w") as f:

        json.dump(self._db, f, indent=2, sort_keys=True)

```

what it does: keeps a JSON database per ingestion root. on each run we check file size/mtime; if they match, we skip. when we do process, we compute and store the SHA256 to make the record deterministic and transportable.

---

## 2) Tiny HDF5 walker (one level down)

**modify**

**willowlab/ingest/file\_index.py**

Replace the HDF5 branch with a one-level walker that samples datasets at root and one-group deep.

# ... keep imports as before ...

```
def _sample_hdf5_obj(obj, budget_items: int = 24, budget_elems: int = 200_000):
```

```
    """
```

Build a small dict of datasets from HDF5 'obj' scanning root and one level deep.

We cap items and elements to avoid RAM blowups.

```
"""

out = {}

def _take(name, dset):

    try:

        if hasattr(dset, "shape"):

            # respect budget

            numel = 1

            for s in getattr(dset, "shape", ()):

                numel *= max(int(s), 1)

            if numel > budget_elems:

                # take a small head slice along first axis (if any)

                sl = (slice(0, min(8, dset.shape[0])),) + tuple(slice(None) for _ in
range(len(dset.shape)-1))

                out[name] = dset[sl]

            else:

                out[name] = dset[()]

        return True

    except Exception:

        return False

    return False


# root datasets

count = 0

for k, v in obj.items():
```

```

if count >= budget_items: break

if hasattr(v, "shape"):
    if _take(k, v): count += 1

# one level deep: groups under root
for k, v in obj.items():
    if count >= budget_items: break
    if hasattr(v, "items"): # group
        for k2, v2 in v.items():
            if count >= budget_items: break
            full = f"{k}/{k2}"
            if hasattr(v2, "shape"):
                if _take(full, v2): count += 1
return out

```

```

def index_willow_dir(data_dir, name_filter=None):
    results = []
    for fn in os.listdir(data_dir):
        if fn.startswith('.'): continue
        if name_filter and name_filter not in fn: continue
        path = os.path.join(data_dir, fn)
        if os.path.isdir(path):
            results.append((path, 'dir', 'unknown', 'directory', None)); continue

```

```

ftype = sniff_file_type(path)

kind, obj = load_any(path)

summary, willow = "", None

if kind == 'array':

    summary = f'array shape={obj.shape} dtype={obj.dtype}'

elif kind == 'mapping':

    keys = list(obj.keys())[:12]

    summary = f'mapping keys={keys}'

    willow = assemble_willow_from_mapping(obj)

elif kind == 'hdf5':

    try:

        sample = _sample_hdf5_obj(obj)

        summary = f'hdf5 sampled keys={list(sample.keys())[:12]}'

        willow = assemble_willow_from_mapping(sample)

    finally:

        try: obj.close()

        except Exception: pass

elif kind == 'text':

    summary = 'text'

else:

    summary = 'unknown/unreadable'

results.append((path, ftype, kind, summary, willow))

```

return results

result: HDF5 files that squirrel data under one extra group get picked up without writing a bespoke loader for every lab's house style.

---

### 3) Merge policy for disagreeing shapes

We'll centralize the merge logic and surface it from the CLI.

**new helper:**

**willowlab/ingest/merge\_policy.py**

```
# willowlab/ingest/merge_policy.py
```

```
import numpy as np
```

```
from typing import Dict, Tuple
```

```
def derive_from_operators(U):
```

```
    """Compute eigenpairs from Floquet operators U."""
```

```
    U = np.asarray(U)
```

```
    if U.ndim == 3:
```

```
        T, N, _ = U.shape
```

```
        evals = np.empty((T, N), dtype=np.complex128)
```

```
        evecs = np.empty_like(U)
```

```
        for t in range(T):
```

```
            w, V = np.linalg.eig(U[t])
```

```
            evals[t] = w; evecs[t] = V
```

```
        return evals, evecs
```

```
    elif U.ndim == 2:
```



```

w, V = np.linalg.eig(U)

return w[None, :], V[None, :, :]

raise ValueError("U must be 2D or 3D array.")

def reconcile_eigenpairs(mapping: Dict, policy: str = "auto") -> Dict:
    """
    Resolve conflicts between supplied eigenpairs and operator-derived ones.

    policy ∈ {'auto', 'prefer_operator', 'prefer_supplied'}
    """

    out = dict(mapping)

    has_U = 'floquet_operators' in out and out['floquet_operators'] is not None and
out['floquet_operators'] != []

    has_ev = 'floquet_eigenvalues' in out and out['floquet_eigenvalues'] is not None and
out['floquet_eigenvalues'] != []

    if not has_U and not has_ev:

        return out # nothing to do

    derived_evals = derived_evecs = None

    if has_U:

        try:

            derived_evals, derived_evecs = derive_from_operators(out['floquet_operators'])

        except Exception:

            derived_evals = derived_evecs = None

```

```
if not has_ev and derived_evals is not None:
```

```
    out['floquet_eigenvalues'] = derived_evals
```

```
    out.setdefault('floquet_eigenvectors', derived_evecs)
```

```
    return out
```

```
if has_ev and derived_evals is None:
```

```
    return out
```

```
# both present → shapes may disagree
```

```
ev_sup = np.asarray(out['floquet_eigenvalues'])
```

```
JT = np.asarray(out.get('JT_scan_points')) if 'JT_scan_points' in out else None
```

```
sup_ok = ev_sup.ndim == 2
```

```
drv_ok = derived_evals is not None and derived_evals.ndim == 2
```

```
# shape sanity
```

```
def _score(ev):
```

```
    if JT is None: return 0
```

```
    return int(ev.shape[0] == JT.shape[0])
```

```
if policy == "prefer_operator" and drv_ok:
```

```
    out['floquet_eigenvalues'] = derived_evals
```

```
    out.setdefault('floquet_eigenvectors', derived_evecs)
```

```
elif policy == "prefer_supplied" and sup_ok:
```

```
    # keep supplied; only fill evecs if missing
```

```
    if 'floquet_eigenvectors' not in out and derived_evecs is not None and ev_sup.shape ==  
    derived_evals.shape:
```

```
        out['floquet_eigenvectors'] = derived_evecs
```

```
    else: # auto
```

```
        # prefer the one that matches JT length; tie-breaker: larger N (more bands)
```

```
    if drv_ok and ( _score(derived_evals) > _score(ev_sup) ):
```

```
        out['floquet_eigenvalues'] = derived_evals
```

```
        out.setdefault('floquet_eigenvectors', derived_evecs)
```

```
    elif drv_ok and sup_ok and _score(derived_evals) == _score(ev_sup):
```

```
        choose_derived = (derived_evals.shape[1] >= ev_sup.shape[1])
```

```
    if choose_derived:
```

```
        out['floquet_eigenvalues'] = derived_evals
```

```
        out.setdefault('floquet_eigenvectors', derived_evecs)
```

```
    # else keep supplied
```

```
    # else keep supplied
```

```
    return out
```

---

## 4) Wire it all into the assembler workflow

### modify

#### willowlab/ingest/assemble\_workflow.py

- add cache usage (skip unchanged),
- include HDF5 walker already via index\_willow\_dir,
- add merge\_policy parameter, and call reconcile\_eigenpairs.

```

# willowlab/ingest/assemble_workflow.py

import os, json, pathlib, glob

from typing import List, Optional

from .sniffers import triage_all_zips

from .file_index import index_willow_dir

from .normalize import to_willow_dataset, persist_dataset

from .merge_policy import reconcile_eigenpairs

from .cache_dir import DirCache

from ..io import load_willow

from ..trinity import WillowTrinityStep1

from ..tests.t_spec_ent import test_duality


PRIORITY_HINTS = ("floquet_braiding", "braiding_2_plaquettes")


def _pick_best(candidates):
    for src, name, w in candidates:
        if ('floquet_eigenvalues' in w) and ('JT_scan_points' in w):
            return (src, name, w)

    for src, name, w in candidates:
        if 'floquet_eigenvalues' in w: return (src, name, w)

    for src, name, w in candidates:
        if 'JT_scan_points' in w: return (src, name, w)

    return (None, None, None)

```

```
def assemble_from_zips(zip_paths: List[str], out_dir: str, expanded_dirs: Optional[List[str]] =
None,
```

```
merge_policy: str = "auto", cache_root: Optional[str] = None):
```

```
    os.makedirs(out_dir, exist_ok=True)
```

```
    cache_root = cache_root or os.path.commonpath(zip_paths + (expanded_dirs or [])) if
zip_paths else (expanded_dirs[0] if expanded_dirs else out_dir)
```

```
    dcache = DirCache(cache_root)
```

```
    # Stage A: zip triage (skip cached zips that haven't changed)
```

```
    zip_paths = [p for p in zip_paths if dcache.should_process(p)]
```

```
    report, assembled = triage_all_zips(zip_paths) if zip_paths else ([], [])
```

```
    # bump priority
```

```
    assembled.sort(key=lambda x: (any(h in x[0].lower() for h in PRIORITY_HINTS) is False))
```

```
    src, inner, mapping = _pick_best(assembled)
```

```
    # Stage B: deep dive expanded dirs (skip cached files internally)
```

```
    if (mapping is None or 'floquet_eigenvalues' not in mapping or 'JT_scan_points' not in
mapping) and expanded_dirs:
```

```
        merged = dict(mapping or {})
```

```
        for d in expanded_dirs:
```

```
            # skip whole dir if none of its files changed
```

```
            inv = index_willow_dir(d)
```

```
            for path, ftype, kind, summary, w in inv:
```

```
                if os.path.isfile(path) and not dcache.should_process(path):
```

```

        continue

    if not w: continue

    for k in
('JT_scan_points','floquet_eigenvalues','floquet_eigenvectors','floquet_operators'):

        if k in w and k not in merged:

            merged[k] = w[k]

    if merged:

        mapping = merged

if mapping is None:

    raise RuntimeError("No usable Willow mapping found in zips or expanded dirs.")

# Merge-policy resolution

mapping = reconcile_eigenpairs(mapping, policy=merge_policy)

meta = {"source_zip": src, "source_entry": inner, "note": f"assembled by WillowLab
(merge_policy={merge_policy})"}

ds = to_willow_dataset(mapping, meta=meta)

npz_path, meta_path = persist_dataset(ds, out_dir)

# Update cache records for processed inputs

for zp in zip_paths:

    dcache.update(zp, bundle=npz_path)

if expanded_dirs:

    for d in expanded_dirs:

```

```

    for leaf in glob.glob(os.path.join(d, "**")):
        if os.path.isfile(leaf):
            dcache.update(leaf, bundle=npz_path)

    dcache.persist()

# Arm the test suite

loaded = load_willow(npz_path)

tri = WillowTrinityStep1(loaded)

inv = tri.compute_all(jt_star=1.0, window=0.05)

summary = {"bundle": npz_path, "meta": meta_path, "invariants": inv}

if (loaded.resolvent_trace is not None) and (loaded.entropy is not None) and
(loaded.effective_energy is not None):

    summary["t_spec_ent"] = test_duality(loaded)

out_json = pathlib.Path(out_dir, "assembly_summary.json")

out_json.write_text(json.dumps(summary, indent=2, default=lambda x: str(x)))

return str(out_json)

```

---

## 5) CLI updates

### modify

#### willowlab/ingest/cli\_assemble.py

Add --merge-policy and --cache-root.

```
# willowlab/ingest/cli_assemble.py
```

```
import argparse, glob
```

```
from .assemble_workflow import assemble_from_zips
```

```
def main():
```

```
    ap = argparse.ArgumentParser(description="Assemble Willow bundles from zips + expanded  
dirs")
```

```
    ap.add_argument("--zip-glob", required=True, help="Glob for zip files (e.g., '/data/*.zip')")
```

```
    ap.add_argument("--expanded-dirs", nargs="*", default=[], help="Dirs expanded next to each  
zip")
```

```
    ap.add_argument("--out", required=True, help="Output directory")
```

```
    ap.add_argument("--merge-policy", choices=["auto", "prefer_operator", "prefer_supplied"],  
default="auto",
```

```
                    help="Resolve eigenpair conflicts")
```

```
    ap.add_argument("--cache-root", default=None, help="Root directory for incremental cache  
DB")
```

```
    args = ap.parse_args()
```

```
    zip_paths = glob.glob(args.zip_glob)
```

```
    summary_path = assemble_from_zips(zip_paths, args.out,  
expanded_dirs=args.expanded_dirs,
```

```
                                merge_policy=args.merge_policy, cache_root=args.cache_root)
```

```
    print(f"Wrote {summary_path}")
```

```
if __name__ == "__main__":
```

```
    main()
```

usage examples:

# 1) first run: full ingest, prefer operator-derived eigenpairs on conflicts

```
python -m willowlab.ingest.cli_assemble \
```



```
--zip-glob "/data/willow/*.zip" \  
  
--expanded-dirs /data/willow/Floquet_Braiding /data/willow/Braiding_2_plaquettes \  
  
--out ./artifacts_assembled \  
  
--merge-policy prefer_operator \  
  
--cache-root /data/willow
```

# 2) subsequent run: unchanged files are skipped instantly via cache

```
python -m willowlab.ingest.cli_assemble \  
  
--zip-glob "/data/willow/*.zip" \  
  
--expanded-dirs /data/willow/Floquet_Braiding /data/willow/Braiding_2_plaquettes \  
  
--out ./artifacts_assembled \  
  
--merge-policy auto \  
  
--cache-root /data/willow
```

---

## sanity checks and edge guards

- cache correctness: we update the cache only after successful assembly. if you want paranoid integrity, switch `should_process` to recompute hashes on every run—but `size+mtime` is a pragmatic default.
- HDF5 walker budget: `budget_items` and `budget_elems` keep memory in check; crank them up if you trust the machine.
- merge-policy auto: aims for the eigenpair set whose time-length matches JT and tie-breaks on band count. that keeps surfaces aligned and avoids silent truncation.

---

you now have an incremental, format-savvy ingestion line that politely skips the déjà vu, yanks eigenpairs from operators when told, and peeks one level deeper into HDF5 labyrinths. point it

at your zips + expanded dirs, choose your merge policy, and let the toolkit auto-populate pristine Willow bundles for the test batteries and cartography engine.